

DATABASE ADMINISTRATION FINAL PROJECT

Exercise 1.1 - Set up the lab environment

- Start the PostgreSQL Server
- Download the lab setup bash file from <https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DB0231EN-SkillsNetwork/labs/Final%20Assignment/postgres-setup.sh>
- Run the bash file

TERMINAL:

- `wget https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DB0231EN-SkillsNetwork/labs/Final%20Assignment/postgres-setup.sh`
- `$ sudo chmod +x postgres-setup.sh`
- `ls -l postgres-setup.sh`
- `./postgres-setup.sh`

Exercise 1.2 - User Management

Task 1.2 - Create a User

- Create a user named backup_operator.

POSTGRE CLI:

- `CREATE USER backup_operator WITH PASSWORD 'backup_operator_password';`

Task 1.3 - Create a Role

- Create a role named backup.

POSTGRE CLI:

- `CREATE ROLE backup;`

Task 1.4 - Grant privileges to the role

- Grant the following privileges to the backup role.
 - `CONNECT ON tolldata DATABASE.`
 - `SELECT ON ALL TABLES IN SCHEMA toll.`

POSTGRE CLI:

- `GRANT CONNECT ON DATABASE tolldata TO backup;`
- `\connect tolldata;`
- `GRANT SELECT ON ALL TABLES IN SCHEMA toll TO backup;`

Task 1.5 - Grant role to an user

- Grant the role backup to backup_operator

POSTGRE CLI:

- `GRANT backup TO backup_operator;`

Exercise 2.1 - Set up the lab environment

- Before you proceed with the assignment, start the MySQL Server.

Exercise 2.2 – Recovery

Task 2.1 - Restore MySQL server using a previous backup

- Download the backup file <https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DB0231EN-SkillsNetwork/labs/Final%20Assignment/billingdata.sql>.
- Restore this file onto MySQL server.
- List the tables in the billing database.

TERMINAL:

```
➤ wget https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DB0231EN-SkillsNetwork/labs/Final%20Assignment/billingdata.sql
```

MySQL CLI:

```
➤ CREATE DATABASE billing;  
➤ USE billing;  
➤ SOURCE billingdata.sql;  
➤ SHOW TABLES;
```

Task 2.2 - Find the table data size

- Find the data size of the table billdata.

MySQL CLI:

```
➤ SELECT table_name, (data_length + index_length)/1024 AS DataSize_kB FROM INFORMATION_SCHEMA.TABLES  
WHERE table_name = 'billdata';
```

Exercise 2.3 – Indexing

Task 2.3 - Baseline query performance

- Write a query to select all rows with a billedamount > 19999 in table billdata.

MySQL CLI:

```
➤ SELECT * FROM billdata WHERE billedamount>19999;  
➤ EXPLAIN SELECT * FROM billdata WHERE billedamount>19999;
```

Task 2.4 - Create an index

- Your customer wants to improve the execution time of the query you wrote in Task 2.3.
- Create an appropriate index to make it run faster.

MySQL CLI:

```
➤ CREATE INDEX billedamount_index ON billdata(billedamount);
```

Task 2.5 - Document the improvement in query performance

- Find out if the index has any impact on query performance.
- Re-run the baseline query of Task 2.3 after creating the index.

MYSQLCLI:

- `SELECT * FROM billdata WHERE billedamount>19999;`
- `EXPLAIN SELECT * FROM billdata WHERE billedamount>19999;`

Exercise 2.4 - Storage Engines

Task 2.6 - Find supported storage engines

- Run a command to find out if your MySQL server supports the MyISAM storage engine.

MYSQL CLI:

- `SHOW ENGINES;`

Task 2.7 - Find the storage engine of a table

- Find the storage engine of the table billdata.

MYSQL CLI:

- `SELECT table_name, engine FROM INFORMATION_SCHEMA.TABLES WHERE table_name = 'billdata';`

Exercise 2.5 - OPTIONAL Exercise (Non-graded) Automation of routine tasks

Bonus Task 2.8 - Write a bash script that performs a backup of all the databases

- mysqldump is a command line tool that performs logical backups of a database.
- Its generic syntax is `mysqldump db_name > backup-file.sql`
- Its extended syntax is `mysqldump --all-databases --user=root --password=NzA4Ny1y > backup-file.sql`
- Write a bash script named mybackup.sh that performs the following tasks.
 - Perform the backup of all databases using the mysqldump
 - Store the output in the file all-databases-backup.sql
 - In the /tmp directory, create a directory named after current date like YYYYMMDD. For example, 20210830
 - Move the file all-databases-backup.sql to /tmp/mysqldumps/<current date>/ directory

MYBACKUP.SH

```
1  #!/bin/sh
2  # The above line tells the interpreter this code needs to be run as a shell script.
3
4  # Set the database name to a variable.
5  DATABASE2='billing'
6
7  # This will be printed on to the screen. In the case of cron job, it will be printed to the logs.
8  echo "Pulling Database: This may take a few minutes"
9
10 # Set the folder where the database backup will be stored
11 backupfolder=/home/theia/backups
12
13 # Number of days to store the backup
14 keep_day=30
15
16 sqlfile=$backupfolder/all-databases-backup.sql
17 zipfile=$backupfolder/all-database-$(date +%d-%m-%Y_%H-%M-%S).gz
18
19 # Create a backup
20
21 if mysqldump $DATABASE2 > $sqlfile ; then
22     echo 'Sql dump created'
23     # Compress backup
24     if gzip -c $sqlfile > $zipfile; then
25         echo 'The backup was successfully compressed'
26     else
27         echo 'Error compressing backupBackup was not created!'
28         exit
29     fi
30     rm $sqlfile
31 else
32     echo 'pg_dump return non-zero code No backup was created!'
33     exit
34 fi
35
36 # Delete old backups
37 find $backupfolder -mtime +$keep_day -delete
38
```

Exercise 3.1 - Restore data

Task 3.1 - Restore the table billing

- Use the billing.csv and restore the CSV file into a table named billing. Write a query to display the first five rows of the table along with the number of rows imported.

QUERY:

➤ `SELECT * FROM billing LIMIT 5;`

Exercise 3.2 - Create a view

Task 3.2 - Create a view named basicbilldetails with the columns customerid, month, billedamount

QUERY:

➤ `CREATE VIEW basicbilldetails AS SELECT customerid, month, billedamount FROM billing;`

Exercise 3.3 - Indexing

Task 3.3 - Baseline query performance

- Write a query to find out all the rows with a billing amount of 19929.
- Hint: Use the command `SELECT strftme('%Y-%m-%d %H:%M:%f', 'now');` before and after your query to display the run time.

QUERY:

- `SELECT strftme('%Y-%m-%d %H:%M:%S:%f', 'now') AS start;`
- `SELECT * FROM billing WHERE billedamount=19929;`
- `SELECT strftme('%Y-%m-%d %H:%M:%S:%f', 'now') AS end;`

Task 3.4 - Create an index

Create an index that can make the query in the previous task faster. Name the index as `billingamount`.

QUERY:

- `CREATE INDEX billingamount ON billing(billedamount);`

Task 3.5 - Document the improvement in query performance

- Find out if the index has any impact on query performance.
- Re-run the query to find out all the rows with a billing amount of 19929.

QUERY:

- `SELECT strftme('%Y-%m-%d %H:%M:%S:%f', 'now') AS start;`
- `SELECT * FROM billing WHERE billedamount=19929;`
- `SELECT strftme('%Y-%m-%d %H:%M:%S:%f', 'now') AS end;`